

# airbnb\_prediction\_model (1)

July 8, 2023

```
[1]: import warnings
warnings.filterwarnings("ignore")
from sklearn.exceptions import FitFailedWarning
```

```
[2]: import pandas as pd
import numpy as np
df = pd.read_csv('train.csv')
df.head(5)
```

```
[2]:      id      name \
0  22267382  Modern and Cozy Large Studio in Brooklyn
1   2473861    Royal Harlem TRIPLEX Home 5 Beds
2  25079703    Sunny East Village Studio
3   9342478    Beautiful, airy, light-filled room
4   4866426    Private Room in Prime Brooklyn Spot
```

```
summary \
0  Modern large studio with new amenities and app...
1  Harlem is back and so gorgeous! Visit and expl...
2  Clean, hip and well designed sun drenched East...
3  Private, spacious, comfortable room in 2-bed f...
4  Comfy, quiet and big private room in a three b...
```

```
space \
0  Our place is a little quiet sanctuary in the h...
1  Harlem is back and so gorgeous! Visit and expl...
2  This is a rare East Village studio with it's h...
3  Big closet, two big windows, tall ceiling and ...
4  This big old apartment that we love and take c...
```

```
description experiences_offered \
0  Modern large studio with new amenities and app...  none
1  Harlem is back and so gorgeous! Visit and expl...  none
2  Clean, hip and well designed sun drenched East...  none
3  Private, spacious, comfortable room in 2-bed f...  none
4  Comfy, quiet and big private room in a three b...  none
```

	neighborhood_overview		notes		transit		access		review_scores_communication	review_scores_location	review_scores_value		instant_bookable	is_business_travel_ready	cancellation_policy		require_guest_profile_picture	require_guest_phone_verification	
0	BAM, Barclays, Brooklyn City Point, Fort Green...		NaN		Subway: 2,3,4,5,A,C,B,Q,G		Washer/Dryer Dishwasher Internet Gym Roof Top ...		10.0	10.0	10.0		t	f	flexible		f	f	
1	HARLEM is a piece of real NY history overflowi...				PUBLIC TRANSPORTATION: Conveniently near all p...		The WHOLE ENTIRE HOUSE ...		9.0	9.0	9.0		t	f	moderate		f	f	
2	East Village is one of the last remaining neig...		NaN				You'll have access to the entire space - it's ...		10.0	10.0	10.0		f	f	moderate		f	f	
3	One block from Morgan L stop. Super cool area...		NaN						NaN	NaN	NaN		f	f	flexible		f	f	
4	I absolutely love this neighborhood - right at...				Super convenient to almost all subway lines. A...		Your room has a very comfortable queen sized b...		10.0	10.0	10.0		f	f	flexible		f	f	

```

4                                     f                                     f

    calculated_host_listings_count  reviews_per_month
0                                1                0.59
1                                3                2.47
2                                1                0.89
3                                1                 NaN
4                                1                3.14

[5 rows x 65 columns]

```

```
[3]: df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 33538 entries, 0 to 33537
Data columns (total 65 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   id                                    33538 non-null  int64
1   name                                33526 non-null  object
2   summary                             32266 non-null  object
3   space                               23038 non-null  object
4   description                         33230 non-null  object
5   experiences_offered                 33538 non-null  object
6   neighborhood_overview               19948 non-null  object
7   notes                               13447 non-null  object
8   transit                             20796 non-null  object
9   access                              19304 non-null  object
10  interaction                         18670 non-null  object
11  house_rules                         19983 non-null  object
12  host_id                             33538 non-null  int64
13  host_name                           33533 non-null  object
14  host_since                          33533 non-null  object
15  host_location                       33424 non-null  object
16  host_about                          20374 non-null  object
17  host_response_time                  17345 non-null  object
18  host_response_rate                   17345 non-null  object
19  host_acceptance_rate                 0 non-null      float64
20  host_is_superhost                   33533 non-null  object
21  host_neighbourhood                  28832 non-null  object
22  host_listings_count                 33533 non-null  float64
23  host_verifications                   33538 non-null  object
24  host_has_profile_pic                 33533 non-null  object
25  host_identity_verified               33533 non-null  object
26  neighbourhood_cleansed               33538 non-null  object
27  neighbourhood_group_cleansed         33538 non-null  object
28  city                                33499 non-null  object

```

29	state	33528	non-null	object
30	zipcode	33053	non-null	object
31	market	33448	non-null	object
32	country_code	33538	non-null	object
33	country	33538	non-null	object
34	property_type	33538	non-null	object
35	room_type	33538	non-null	object
36	accommodates	33538	non-null	int64
37	bathrooms	33478	non-null	float64
38	bedrooms	33505	non-null	float64
39	beds	33507	non-null	float64
40	bed_type	33538	non-null	object
41	amenities	33538	non-null	object
42	square_feet	341	non-null	float64
43	price	33538	non-null	int64
44	guests_included	33538	non-null	int64
45	extra_people	33538	non-null	object
46	minimum_nights	33538	non-null	int64
47	maximum_nights	33538	non-null	int64
48	number_of_reviews	33538	non-null	int64
49	first_review	26591	non-null	object
50	last_review	26593	non-null	object
51	review_scores_rating	25874	non-null	float64
52	review_scores_accuracy	25844	non-null	float64
53	review_scores_cleanliness	25859	non-null	float64
54	review_scores_checkin	25829	non-null	float64
55	review_scores_communication	25849	non-null	float64
56	review_scores_location	25830	non-null	float64
57	review_scores_value	25827	non-null	float64
58	instant_bookable	33538	non-null	object
59	is_business_travel_ready	33538	non-null	object
60	cancellation_policy	33538	non-null	object
61	require_guest_profile_picture	33538	non-null	object
62	require_guest_phone_verification	33538	non-null	object
63	calculated_host_listings_count	33538	non-null	int64
64	reviews_per_month	26591	non-null	float64

dtypes: float64(14), int64(9), object(42)  
memory usage: 16.6+ MB

```
[4]: def convert_tf(x):
      if x == 't':
          return 1
      elif x == 'f':
          return 0
      else:
          pass
      df['host_is_superhost'] = df['host_is_superhost'].apply(convert_tf)
```

```
[5]: df['host_is_superhost']
```

```
[5]: 0      0.0
      1      0.0
      2      0.0
      3      0.0
      4      1.0
      ...
      33533  0.0
      33534  0.0
      33535  0.0
      33536  0.0
      33537  0.0
      Name: host_is_superhost, Length: 33538, dtype: float64
```

```
[6]: df['market'].value_counts()
```

```
[6]: New York      33425
      Other (Domestic)  10
      Catskills and Hudson Valley  2
      Los Angeles  2
      Adirondacks  2
      Boston  1
      San Francisco  1
      Jamaica South Coast  1
      New Orleans  1
      Agra  1
      Paris  1
      Kyoto  1
      Name: market, dtype: int64
```

```
[7]: df.select_dtypes(include=['int', 'float']).drop(columns=['id', 'host_id'])
```

```
[7]:      host_acceptance_rate  host_is_superhost  host_listings_count  \
0      NaN      0.0      1.0
1      NaN      0.0      4.0
2      NaN      0.0      1.0
3      NaN      0.0      1.0
4      NaN      1.0      1.0
...      ...      ...      ...
33533      NaN      0.0      1.0
33534      NaN      0.0      3.0
33535      NaN      0.0      1.0
33536      NaN      0.0      4.0
33537      NaN      0.0      1.0

      accommodates  bathrooms  bedrooms  beds  square_feet  price  \
```

0	2	1.0	1.0	1.0	NaN	145
1	8	1.0	3.0	5.0	NaN	175
2	2	1.0	0.0	1.0	NaN	180
3	1	1.0	1.0	1.0	NaN	42
4	2	1.0	1.0	1.0	NaN	80
...	...	...	...	...	...	...
33533	2	1.0	1.0	1.0	NaN	110
33534	8	3.0	4.0	4.0	NaN	1195
33535	1	NaN	1.0	1.0	NaN	50
33536	2	1.0	1.0	1.0	NaN	60
33537	8	2.0	3.0	4.0	NaN	325

	guests_included	...	number_of_reviews	review_scores_rating	\
0	2	...	6	100.0	
1	3	...	137	91.0	
2	1	...	3	100.0	
3	1	...	0	NaN	
4	1	...	144	97.0	
...	...	...	...	...	
33533	1	...	19	87.0	
33534	1	...	86	95.0	
33535	1	...	0	NaN	
33536	1	...	1	100.0	
33537	4	...	56	95.0	

	review_scores_accuracy	review_scores_cleanliness	\
0	10.0	10.0	
1	9.0	9.0	
2	10.0	9.0	
3	NaN	NaN	
4	10.0	10.0	
...	...	...	
33533	9.0	8.0	
33534	10.0	9.0	
33535	NaN	NaN	
33536	10.0	10.0	
33537	9.0	9.0	

	review_scores_checkin	review_scores_communication	\
0	10.0	10.0	
1	9.0	9.0	
2	9.0	10.0	
3	NaN	NaN	
4	10.0	10.0	
...	...	...	
33533	9.0	9.0	
33534	10.0	10.0	

33535	NaN	NaN
33536	10.0	10.0
33537	10.0	10.0

	review_scores_location	review_scores_value \
0	10.0	10.0
1	9.0	9.0
2	10.0	10.0
3	NaN	NaN
4	10.0	10.0
...	...	...
33533	9.0	8.0
33534	10.0	9.0
33535	NaN	NaN
33536	10.0	10.0
33537	10.0	9.0

	calculated_host_listings_count	reviews_per_month
0	1	0.59
1	3	2.47
2	1	0.89
3	1	NaN
4	1	3.14
...	...	...
33533	1	0.44
33534	1	2.28
33535	1	NaN
33536	3	0.10
33537	1	2.76

[33538 rows x 22 columns]

```
[8]: baseline_features = ['accommodates',
                          'bathrooms',
                          'bedrooms',
                          'beds',
                          'guests_included',
                          'number_of_reviews',
                          'reviews_per_month',
                          'neighbourhood_group_cleansed',
                          'room_type',
                          'host_is_superhost',
                          'property_type'
                          ]
```


```
[9]: import pandas as pd
import seaborn as sns
```

```
import matplotlib.pyplot as plt

# create a sample dataframe
plt.figure(figsize=(20,20))
# calculate the correlation matrix
corr_matrix = df.select_dtypes(include=['int', 'float']).drop(columns=['id', 'host_id']).corr()

# plot the correlation heatmap
sns.heatmap(corr_matrix, annot=True)


# show the plot
plt.show()
```





```
[10]: df['price'].hist(bins = 50)
```

```
[10]: <AxesSubplot:>
```



```
[11]: considering_features = ['host_is_superhost', 'state', 'country',
    ↪ 'property_type', 'accommodates', 'bedrooms', 'number_of_reviews',
    ↪ 'review_scores_rating']
```

```
[12]: basic_feature = ['host_is_superhost', 'state', 'bedrooms',
    ↪ 'number_of_reviews', 'review_scores_rating', 'accommodates']
```

```
for bf in basic_feature:
    print(bf + ": ", df[bf].unique(), df[bf].dtype)
```

```
host_is_superhost: [ 0.  1. nan] float64
state: ['NY' 'Ny' 'New York' nan 'MP' 'NJ' 'ny' 'CA'] object
bedrooms: [ 1.  3.  0.  2.  5.  4. nan  6.  9.  7.  8. 10.] float64
number_of_reviews: [ 6 137  3  0 144 12  1 30  2 10 118 173 16  8 66
29 171 108
157 47 27 19 127 24  7 25 70 31 17 107 11 41 102  4 104 84
85 134  5 76 73  9 167 22 83 72 313 14 35 23 45 153 142 68
20 208 91 13 26 50 42 88 21 92 98 18 89 52 40 36 103 59
15 176 28 44 55 53 106 32 43 57 64 58 69 213 38 51 205 156
```


```

121 113 123 78 56 39 49 95 128 146 115 65 135 166 34 37 75 241
97 67 162 63 226 120 93 125 90 33 71 177 112 74 151 129 141 109
204 86 99 187 277 193 150 139 184 138 81 296 110 100 46 170 278 212
119 54 96 169 229 87 62 246 82 60 131 61 333 161 154 230 168 221
209 111 190 197 207 80 114 155 165 132 158 48 305 195 174 116 101 200
94 317 202 77 79 244 273 206 201 194 266 242 140 149 240 243 186 303
145 311 126 191 147 269 117 260 124 164 231 122 188 233 245 409 557 214
133 105 216 315 143 136 160 319 175 182 282 308 180 159 222 224 210 181
284 529 172 152 163 215 268 239 196 218 203 397 263 290 198 265 347 217
178 148 276 130 473 225 179 256 228 270 236 237 238 234 227 192 283 348
306 338 261 232 247 264 267 250 211 354 356 185 199 359 301 418 279 343
255 189 259 275 300 324 326 455 272 183 421 219 334 312 410 262 288 391
375 293 341 358 252 251 258 285 253 220 294 297 405 362 463 287 307 394
304 412 401 235 249 248 309 393 286 482 378 289] int64
review_scores_rating: [100. 91. nan 97. 93. 96. 80. 90. 73. 98. 92.
95. 94. 99.
84. 89. 88. 83. 85. 60. 75. 86. 20. 50. 69. 87. 40. 79.
76. 70. 82. 74. 78. 65. 77. 45. 64. 67. 81. 57. 72. 53.
66. 71. 58. 55. 68. 62. 47. 30.] float64
accommodates: [ 2  8  1  6  4 16  5 12  3  7  9 10 13 11 14 15] int64

```


```
[13]: df['host_is_superhost'].value_counts().plot(kind='bar')
```

```
[13]: <AxesSubplot:>
```




```
[14]: df['bedrooms'].value_counts().plot(kind='bar')
```

```
[14]: <AxesSubplot:>
```




```
[15]: df['number_of_reviews'].value_counts().plot(kind='hist', ec= 'w')
```

```
[15]: <AxesSubplot:ylabel='Frequency'>
```




```
[16]: df['review_scores_rating'].value_counts().plot(kind='hist', ec= 'w')
```

```
[16]: <AxesSubplot:ylabel='Frequency'>
```



```
[17]: df['accommodates'].value_counts().plot(kind='bar')
```

```
[17]: <AxesSubplot:>
```



```
[18]: df.select_dtypes(include=['int', 'float']).drop(columns=['id', 'host_id'])
```

```
[18]:
```

	host_acceptance_rate	host_is_superhost	host_listings_count	\
0	NaN	0.0	1.0	
1	NaN	0.0	4.0	
2	NaN	0.0	1.0	
3	NaN	0.0	1.0	
4	NaN	1.0	1.0	
...	...	...	...	
33533	NaN	0.0	1.0	
33534	NaN	0.0	3.0	
33535	NaN	0.0	1.0	
33536	NaN	0.0	4.0	
33537	NaN	0.0	1.0	

	accommodates	bathrooms	bedrooms	beds	square_feet	price	\
0	2	1.0	1.0	1.0	NaN	145	
1	8	1.0	3.0	5.0	NaN	175	
2	2	1.0	0.0	1.0	NaN	180	
3	1	1.0	1.0	1.0	NaN	42	
4	2	1.0	1.0	1.0	NaN	80	

...	...	...	...	...	...	...
33533	2	1.0	1.0	1.0	NaN	110
33534	8	3.0	4.0	4.0	NaN	1195
33535	1	NaN	1.0	1.0	NaN	50
33536	2	1.0	1.0	1.0	NaN	60
33537	8	2.0	3.0	4.0	NaN	325

	guests_included	...	number_of_reviews	review_scores_rating	\
0	2	...	6	100.0	
1	3	...	137	91.0	
2	1	...	3	100.0	
3	1	...	0	NaN	
4	1	...	144	97.0	
...	...	...	...	...	
33533	1	...	19	87.0	
33534	1	...	86	95.0	
33535	1	...	0	NaN	
33536	1	...	1	100.0	
33537	4	...	56	95.0	

	review_scores_accuracy	review_scores_cleanliness	\
0	10.0	10.0	
1	9.0	9.0	
2	10.0	9.0	
3	NaN	NaN	
4	10.0	10.0	
...	...	...	
33533	9.0	8.0	
33534	10.0	9.0	
33535	NaN	NaN	
33536	10.0	10.0	
33537	9.0	9.0	

	review_scores_checkin	review_scores_communication	\
0	10.0	10.0	
1	9.0	9.0	
2	9.0	10.0	
3	NaN	NaN	
4	10.0	10.0	
...	...	...	
33533	9.0	9.0	
33534	10.0	10.0	
33535	NaN	NaN	
33536	10.0	10.0	
33537	10.0	10.0	

	review_scores_location	review_scores_value	\
--	------------------------	---------------------	---

0	10.0	10.0
1	9.0	9.0
2	10.0	10.0
3	NaN	NaN
4	10.0	10.0
...	...	...
33533	9.0	8.0
33534	10.0	9.0
33535	NaN	NaN
33536	10.0	10.0
33537	10.0	9.0

	calculated_host_listings_count	reviews_per_month
0	1	0.59
1	3	2.47
2	1	0.89
3	1	NaN
4	1	3.14
...	...	...
33533	1	0.44
33534	1	2.28
33535	1	NaN
33536	3	0.10
33537	1	2.76

[33538 rows x 22 columns]

```
[19]: missing = df.isna().sum()

missing /= df.shape[0]
missing *=100
missing = missing.to_frame().rename(columns={0:'Precent Of Missing Values'})
missing
```

```
[19]:
```

	Precent Of Missing Values
id	0.000000
name	0.035780
summary	3.792713
space	31.307770
description	0.918361
...	...
cancellation_policy	0.000000
require_guest_profile_picture	0.000000
require_guest_phone_verification	0.000000
calculated_host_listings_count	0.000000
reviews_per_month	20.713817

[65 rows x 1 columns]

```
[20]: missing.loc[baseline_features]
```

```
[20]:
```

	Precent Of Missing Values
accommodates	0.000000
bathrooms	0.178902
bedrooms	0.098396
beds	0.092432
guests_included	0.000000
number_of_reviews	0.000000
reviews_per_month	20.713817
neighbourhood_group_cleansed	0.000000
room_type	0.000000
host_is_superhost	0.014908
property_type	0.000000

```
[21]: from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.impute import SimpleImputer
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.datasets import fetch_openml
from sklearn.model_selection import train_test_split
from xgboost import XGBRegressor
from sklearn.metrics import mean_squared_error
from sklearn.model_selection import RandomizedSearchCV
```

### 0.0.1 XGB

#### Search Best param for XGBoost

```
[22]: # df['room_type']=df['room_type'].str.split('/', expand=True)[0].str.split('_',
↳ ', expand=True)[0]
# df['room_type'].value_counts()
```

```
[96]: from sklearn.base import BaseEstimator, TransformerMixin
from sklearn.impute import SimpleImputer
from sklearn.preprocessing import StandardScaler

class Impute_Standardize(BaseEstimator, TransformerMixin):
    def __init__(self, strategy='mean', fill_value=None):
        self.strategy = strategy
        self.fill_value = fill_value
        self.imputer = SimpleImputer(strategy=self.strategy,
                                      fill_value=self.fill_value)
        self.scaler = StandardScaler()
```



```

def fit(self, X, y=None):
    self.imputer.fit(X)
    self.scaler.fit(X)
    return self

def transform(self, X, y=None):
    X_imputed = self.imputer.transform(X)
    X_imputed = pd.DataFrame(X_imputed, columns=self.imputer.
↪get_feature_names_out())
    X_scaled = self.scaler.transform(X_imputed)
    return X_scaled

```

```

[93]: preprocessor = ColumnTransformer([('ohe', OneHotEncoder(),
                                       ('neighbourhood_group_cleansed',
↪'room_type', 'property_type', 'host_is_superhost')),

                                       ('imp-mode',
↪SimpleImputer(strategy='most_frequent'),
                                       ('bathrooms', 'bedrooms', 'beds')),

                                       ('imp-mean', Impute_Standardize(strategy='mean'),
                                       ('review_scores_accuracy')),

                                       ('imp-cons', SimpleImputer(strategy='constant',
↪fill_value=0),
                                       ('reviews_per_month')),

                                       ('std', StandardScaler(), ('accommodates',
                                       'guests_included',
                                       'number_of_reviews'))])

```

```

[91]: baseline_features = ['accommodates',
                           'bathrooms',
                           'bedrooms',
                           'beds',
                           'guests_included',
                           'number_of_reviews',
                           'reviews_per_month',
                           'neighbourhood_group_cleansed',
                           'room_type',
                           'host_is_superhost',
                           'property_type',
                           'review_scores_accuracy'
                           ]

```

```
[95]: one_hot_features = ['neighbourhood_group_cleansed',
                        'room_type',
                        'property_type',
                        'host_is_superhost',
                        'bathrooms',
                        'bedrooms',
                        'beds']
```

```
impute_standardize_mean = ['review_scores_accuracy']
```

```
simple_imputer = ['reviews_per_month']
```

```
standard_scaler = ['accommodates',
                  'guests_included',
                  'number_of_reviews']
```

```
[105]: df_x, df_y = df.drop(columns=['price']), df['price']
```

```
one_hot_features = ['neighbourhood_group_cleansed', 'room_type',
                    'property_type',
                    'host_is_superhost', 'bathrooms', 'bedrooms', 'beds']
```

```
impute_standardize_mean = ['review_scores_accuracy']
```

```
simple_imputer = ['reviews_per_month']
```

```
standard_scaler = ['accommodates', 'guests_included', 'number_of_reviews']
```

```
one_hot_transformer = Pipeline(steps=[
    ('imputer', SimpleImputer(strategy='most_frequent')),
    ('onehot', OneHotEncoder(handle_unknown='ignore', sparse=False))
])
```

```
impute_standardize_mean_transformer = Pipeline(steps=[
    ('imputer', Impute_Standardize(strategy='mean'))
])
```

```
simple_imputer_transformer = Pipeline(steps=[
    ('imputer', SimpleImputer(strategy='constant'))
])
```

```
standard_scaler_transformer = Pipeline(steps=[
    ('scaler', StandardScaler())
])
```

```
preprocessor = ColumnTransformer(transformers=[
    ('ohe', one_hot_transformer, one_hot_features),
```

```

    ('ism', impute_standardize_mean_transformer, impute_standardize_mean),
    ('si', simple_imputer_transformer, simple_imputer),
    ('ss', standard_scaler_transformer, standard_scaler)
])

model = Pipeline([
    ('preprocessor', preprocessor),
    ('classifier', XGBRegressor())
])

X_train, X_test, y_train, y_test = train_test_split(df_x, df_y, test_size=0.2,
    random_state=42)

param_grid = {
    'classifier__learning_rate': [0.01, 0.05, 0.1, 0.2],
    'classifier__n_estimators': range(50, 200, 10),
    'classifier__max_depth': range(3, 10),
    'classifier__colsample_bytree': [0.5, 0.7, 1],
    'classifier__gamma': [0, .25, 1.0],
    'classifier__subsample': [0.5, 0.7, 1],
    'classifier__reg_lambda': [0, 1.0, 10.0]
}

random_search = RandomizedSearchCV(
    model,
    param_distributions=param_grid,
    n_iter=10,
    cv=5,
    n_jobs=-1
)

# random_search.fit(X_train, y_train)

# print('Best parameters:', random_search.best_params_)
# print('Best score:', -random_search.best_score_)

```

```

[114]: from sklearn.model_selection import RandomizedSearchCV, train_test_split
from sklearn.pipeline import Pipeline, make_pipeline
from sklearn.compose import ColumnTransformer
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.impute import SimpleImputer
from xgboost import XGBRegressor
import pandas as pd
import numpy as np

one_hot_features = ['neighbourhood_group_cleansed',

```

```

        'room_type',
        'property_type',
        'host_is_superhost',
        'bathrooms',
        'bedrooms',
        'beds']

impute_standardize_mean = ['review_scores_accuracy']

simple_imputer = ['reviews_per_month']

standard_scaler = ['accommodates',
                  'guests_included',
                  'number_of_reviews']

df_x, df_y = df.drop(columns=['price']), df['price']

one_hot_transformer = Pipeline(steps=[
    ('imputer', SimpleImputer(strategy='most_frequent')),
    ('onehot', OneHotEncoder(handle_unknown='ignore'))
])

impute_standardize_mean_transformer = Pipeline(steps=[
    ('imputer', SimpleImputer(strategy='most_frequent'))
])

simple_imputer_transformer = Pipeline(steps=[
    ('imputer', SimpleImputer(strategy='constant'))
])

standard_scaler_transformer = Pipeline(steps=[
    ('simple_imputer', SimpleImputer(strategy='most_frequent')),
    ('imputer', StandardScaler())
])

preprocessor = ColumnTransformer(transformers=[
    ('ohe', one_hot_transformer, one_hot_features),
    ('ism', impute_standardize_mean_transformer, impute_standardize_mean),
    ('si', simple_imputer_transformer, simple_imputer),
    ('ss', standard_scaler_transformer, standard_scaler)
])

# preprocessor.fit_transform(df)

model = Pipeline([
    ('preprocessor', preprocessor),
    ('classifier', XGBRegressor())
])

```

```

])

X_train, X_test, y_train, y_test = train_test_split(df_x, df_y, test_size=0.2,
↳random_state=42)

param_grid = {
    'classifier__learning_rate': [0.01, 0.05, 0.1, 0.2],
    'classifier__n_estimators': range(50, 200, 10),
    'classifier__max_depth': range(3, 10),
    'classifier__colsample_bytree': [0.5, 0.7, 1],
    'classifier__gamma': [0, .25, 1.0],
    'classifier__subsample': [0.5, 0.7, 1],
    'classifier__reg_lambda': [0, 1.0, 10.0]
}

random_search = RandomizedSearchCV(
    model,
    param_distributions=param_grid,
    n_iter=10,
    cv=5,
    n_jobs=-1
)

random_search.fit(X_train, y_train)

print('Best parameters:', random_search.best_params_)

```

```

Best parameters: {'classifier__subsample': 1, 'classifier__reg_lambda': 1.0,
'classifier__n_estimators': 120, 'classifier__max_depth': 8,
'classifier__learning_rate': 0.05, 'classifier__gamma': 0.25,
'classifier__colsample_bytree': 0.5}

```

```

[115]: # best_xgb_model = Pipeline([
#       ('preproc', preprocessor),
#       ('classifier', XGBRegressor(
#           subsample= 0.7,
#           reg_lambda= 1.0,
#           n_estimators= 80,
#           max_depth= 5,
#           learning_rate = 0.05,
#           gamma = 0.25,
#           colsample_bytree=1)
#       )
# ])

```

```

[116]: best_xgb_model = random_search.best_estimator_

```

```
[117]: df_x, df_y = df.drop(columns=['price']), df['price']

# split the data into training and testing sets
xgb_X_train, xgb_X_test, xgb_y_train, xgb_y_test = train_test_split(df_x, df_y,
    ↪test_size=0.2, random_state=42)

# Fit the pipeline to the training data
best_xgb_model.fit(xgb_X_train, xgb_y_train)

xgb_y_pred_train = best_xgb_model.predict(xgb_X_train)
print('train_RMSE', mean_squared_error(xgb_y_pred_train, xgb_y_train)**0.5)

xgb_y_pred_test = best_xgb_model.predict(xgb_X_test)
print('test_RMSE', mean_squared_error(xgb_y_pred_test, xgb_y_test)**0.5)
```

```
train_RMSE 78.69322389542123
test_RMSE 95.10140155751522
```

```
[118]: best_xgb_model.fit(df_x,df_y)

xgb_y_pred = best_xgb_model.predict(df_x)
print('train_RMSE', mean_squared_error(xgb_y_pred, df_y)**0.5)
```

```
train_RMSE 80.2805870671499
```

### Apply on Test set

```
[119]: test = pd.read_csv('test.csv')
test.head()
```

```
[119]:      id                                name \
0  19307997          Super Lux 2BR in Downtown Manhattan
1  20176193    Vintage Eclectic Brownstone Pad in Brooklyn
2  19485371          Spacious Harlem Hideaway
3  13079990    Spacious private room in Brooklyn
4  22339757  *Dg) Delightful Private Room 20 min to Manhattan
```

```
summary \
0  Prepare to be WOWED! This spectacularly bright...
1  Ideal for romantic, creative types, this is an...
2  Postive Vibes . This is our Harlem tree house,...
3  Newly renovated apartment, its a 3 bedroom apa...
4  Hi my home is only 2 blocks from the subway st...
```

```
space \
0  Top of the line Wolf and Sub-Zero appliances, ...
1  Not your typical New York abode, my apartment ...
```

2 The private room is very spacious and cozy. Th...  
 3 3 bedroom apartment, 1 full bathroom, living r...  
 4 NaN

	description	experiences_offered \
0	Prepare to be WOWED! This spectacularly bright...	none
1	Ideal for romantic, creative types, this is an...	none
2	Postive Vibes . This is our Harlem tree house,...	none
3	Newly renovated apartment, its a 3 bedroom apa...	none
4	Hi my home is only 2 blocks from the subway st...	none

	neighborhood_overview \
0	NaN
1	Bed Stuy is a diverse historic neighborhood wi...
2	You are in a Cultural Haven full of restaurant...
3	NaN
4	NaN

	notes \
0	NaN
1	This is an actual unique living experience whe...
2	We also keep cucumber water in the fridge feel...
3	NaN
4	NaN

	transit \
0	NaN
1	Close to buses and subways there is also free ...
2	Train, uber or a taxi. (Extremely taxi accessi...
3	There is the Mta 3 train Sutter stop, also the...
4	NaN

	access ... \
0	NaN ...
1	Entrance hallway, living room, bedroom, kitche... ..
2	Private Room, Kitchen And Bathroom ...
3	NaN ...
4	NaN ...

	review_scores_communication	review_scores_location	review_scores_value \
0	NaN	NaN	NaN
1	10.0	10.0	10.0
2	10.0	10.0	10.0
3	9.0	8.0	9.0
4	8.0	8.0	8.0

instant_bookable	is_business_travel_ready	cancellation_policy \
------------------	--------------------------	-----------------------

0	f	f	flexible
1	f	f	flexible
2	t	f	flexible
3	f	f	flexible
4	t	f	strict_14_with_grace_period

	require_guest_profile_picture	require_guest_phone_verification	\
0	f	f	
1	f	f	
2	f	f	
3	f	f	
4	f	f	

	calculated_host_listings_count	reviews_per_month
0	1	NaN
1	1	1.48
2	1	0.37
3	1	0.23
4	9	1.53

[5 rows x 64 columns]

```
[120]: xgb_y_pred = best_xgb_model.predict(df_x)
print('train_RMSE', mean_squared_error(xgb_y_pred, df_y)**0.5)
```

train\_RMSE 80.2805870671499

```
[121]: xgb_df_predict_y = best_xgb_model.predict(test)
output_df = pd.DataFrame()
output_df['Id'] = test['id']
output_df['Predicted'] = xgb_df_predict_y
output_df.to_csv('xgb_prediction.csv', index = False)
```

## 0.0.2 RandomForest

```
[153]: from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.impute import SimpleImputer
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.datasets import fetch_openml
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import mean_squared_error
```

Search Best Param\_\_



```
[158]: # df_x, df_y = df.drop(columns=['price']), df['price']

# categorical_transformer = Pipeline(steps=[
#     ('imputer', SimpleImputer(strategy='most_frequent')),
#     ('onehot', OneHotEncoder(handle_unknown='ignore'))
# ])

# preprocessor = ColumnTransformer(transformers=[
#     ('ohe', categorical_transformer, baseline_features)
# ])

# rf_model = Pipeline([
#     ('preproc', preprocessor),
#     ('classifier', RandomForestRegressor())
# ])

# param_grid = {
#     'classifier__n_estimators': [10, 50, 100],
#     'classifier__max_depth': [None, 10, 20],
# }

# random_search = RandomizedSearchCV(
#     rf_model,
#     param_distributions=param_grid,
#     n_iter=10,
#     cv=5,
#     n_jobs=-1,
#     random_state=42
# )

# random_search.fit(df_x, df_y)

# print(random_search.best_params_)
```

```
[154]: df_x, df_y = df.drop(columns=['price']), df['price']

categorical_transformer = Pipeline(steps=[
    ('imputer', SimpleImputer(strategy='most_frequent')),
    ('onehot', OneHotEncoder(handle_unknown='ignore'))
])

preprocessor = ColumnTransformer(transformers=[
    ('ohe', categorical_transformer, baseline_features)
])

rf_model = Pipeline([
    ('preproc', preprocessor),
```

```

        ('classifier', RandomForestRegressor(n_estimators=50,max_depth=10))
    ])

    # split the data into training and testing sets
    rf_X_train, rf_X_test, rf_y_train, rf_y_test = train_test_split(df_x, df_y,
        ↪test_size=0.2, random_state=42)

    # Fit the pipeline to the training data
    rf_model.fit(rf_X_train, y_train)

```

```

[154]: Pipeline(steps=[('preproc',
                        ColumnTransformer(transformers=[('ohe',
                                                         Pipeline(steps=[('imputer',
                                                         SimpleImputer(strategy='most_frequent')),
                                                         ('onehot',
                                                         OneHotEncoder(handle_unknown='ignore'))])),
                        ['accommodates', 'bathrooms',
                        'bedrooms', 'beds',
                        'guests_included',
                        'number_of_reviews',
                        'reviews_per_month',
                        'neighbourhood_group_cleansed',
                        'room_type']]])),
        ('classifier',
         RandomForestRegressor(max_depth=10, n_estimators=50)))]

```

```

[157]: rf_y_pred_train = rf_model.predict(rf_X_train)
        print('train_RMSE', mean_squared_error(rf_y_pred_train, rf_y_train)**0.5)

        rf_y_pred_test = rf_model.predict(rf_X_test)
        print('test_RMSE', mean_squared_error(rf_y_pred_test, rf_y_test)**0.5)

```

```

train_RMSE 84.37952399596651
test_RMSE 99.68500273229499

```

```
[ ]:
```